
Implementation of a hearing-impairment simulator based on TEENSY

Special Course

Supervisor:

Bastian Epp

bepp@elektro.dtu.dk

Author:

Vassili Cruchet, s172226

vassili.cruchet@gmail.com

DTU Electrical Engineering
Technical University of Denmark
Spring, 2018

Contents

1	Introduction	1
2	Frequency smearing theory	1
2.1	Basic principle	1
2.2	Algorithm implementation	2
3	<i>MATLAB</i> implementation	4
3.1	Functions description	4
3.1.1	A_s = calc_smear_matrix(fs, N, b)	4
3.1.2	A = calc_audit_filt(fs, N, b)	4
3.1.3	Y = smearing(X, A_s)	5
3.2	Test scripts	5
3.2.1	test_frequency_smearing.m	5
3.2.2	frequency_smearing.m	6
3.3	Comparison with Bear and Moore results [2]	7
3.4	Scripts and functions used for C-code generation	9
3.4.1	Row-indexed sparse storage	9
3.4.2	generate_smear_matrix.m	9
4	Implementation on <i>TEENSY</i> board	11
4.1	Hardware description	11
4.2	Software description	11
4.2.1	Real-time block processing	12
4.2.2	Smearing algorithm	13
4.3	Functions description	13

4.3.1	read_array_form_queue	13
4.3.2	read_frame_from_buffer	14
4.3.3	overlap_add	14
4.3.4	smearing_uncomp and smearing_comp	14
4.3.5	spr sax	15
4.3.6	neg_freq	15
4.3.7	fft_init	16
4.3.8	set_periph	16
4.3.9	setI2SFreq	17
4.3.10	read_button	17
4.3.11	create_hann_window and create_sqrthann_window	17
5	Measurements on TEENSY	18
5.1	Experimental setup	18
5.2	Results and discussion	18
6	Limitations, improvements and future work	21
6.1	Processing time	21
6.2	Memory usage, operation flexibility and characterization	21
6.3	Hardware improvement	22
6.4	Algorithm improvement	22
6.5	Future work	22
	References	23

1 Introduction

Baer and Moore [2] described that some effects of hearing impairment can be modeled by frequency smearing. They tested an algorithm on pre-recorded audio samples and were able to show a correlation between speech intelligibility and frequency smearing, especially in presence of noise. The idea and goal of this project is to implement this algorithm on an embedded platform to process audio in real-time. This could then be used as a demonstrator in different events where people could wear headphones and experience this aspect of hearing impairment and the deficit in communication that comes with it.

In practice, the algorithm was first implemented in MATLAB to test it off-line with different parameters. Secondly, the algorithm was implemented on a microcontroller-based platform (TEENSY 3.6) which integrate a DSP co-processor with which complete libraries can be used.

The whole code and documentation can be found on a *git-hub* repository ¹.

2 Frequency smearing theory

2.1 Basic principle

As explained in [4], the ear identifies frequencies using a frequency-to-place transformation in the cochlea that could be modeled as a bank of asymmetric band-pass filters. The width of these so called auditory filters depends on the center frequency and can be represented using a *roex*(p) function.

$$W(g) = (1 + pg)e^{-pg}, \quad g = \frac{|f - f_0|}{f_0}, \quad (1)$$

where g is the normalized frequency, f_0 the center frequency and p determines the tuning of the filter. Note that this model assumes a symmetric auditory filter. Moreover, one can define the equivalent rectangular bandwidth (ERB) as follows [4] :

$$\text{ERB}_{\text{roex}} = \frac{4f_0}{p} \simeq 24.7(0.000437f_0 + 1) \quad (2)$$

\uparrow
 normal hearing

¹https://github.com/cruchet/teensy_FrequencySmearing




Figure 1: Auditory filter centered on $f_0 = 1kHz$

Equation 2 can then be used to estimate the parameter p and frequency smearing can be simulated by broadening of the auditory filter. This is done by dividing p by a broadening factor b . Figure 1 shows auditory filters for different values of b .

2.2 Algorithm implementation

The broadening explained above is accomplished using a matrix multiplication. The power spectrum is represented in a vector where the n -th component is the power at n times the frequency resolution. When this vector is multiplied with the smearing matrix as in equation 3, this component is replaced by the convolution of itself with the broadened auditory filter centered on the corresponding frequency.

$$Y(n) = \sum_i A_s(n, i)X(i) \iff \mathbf{Y} = \mathbf{A}_s \mathbf{X}, \quad (3)$$

where \mathbf{X} and \mathbf{Y} are the input and output spectrum vectors and \mathbf{A}_s is the matrix containing the normalized smeared auditory filter. $A(n, i)$ is then the value of the filter centered on the n -th frequency at the i -th frequency. It is calculated as follow:

$$\mathbf{A}_s = \mathbf{A}_N^{-1} \mathbf{A}_W, \quad (4)$$

with

$$A_N(n, i) = \frac{1}{ERB}(1 + pg)e^{-pg}, \text{ and } A_W(n, i) = (1 + \frac{pg}{b})e^{-\frac{pg}{b}}, \quad (5)$$

using the normalized frequency

$$g(n, i) = \frac{|f_i - f_n|}{f_n}. \quad (6)$$

As seen in equation 4, it is necessary to multiply the widened filter \mathbf{A}_W with the inverse of the normal auditory filter \mathbf{A}_N as well as to divide it with by the normal hearing ERB to compensate for the listener's own auditory filter.

Figure 2 shows the signal path when processed to be smeared. The window functions used before FFT and after IFFT are square root of hann window. The square root is needed as windowing occurs to times. This will achieve exact reconstruction with an overlap of 50%. As seen on figure 3, the smearing only affects the magnitude and the original phase is kept to preserve the time domain waveform, as suggested by Baer and Moore in [2].



Figure 2: Diagram of the operations to perform frequency smearing. The time-domain frame is first windowed with a square-root of hann window. After FFT, the power spectrum is processed and converted back to time-domain. It is windowed again before overlap-and-add with 50% overlap.

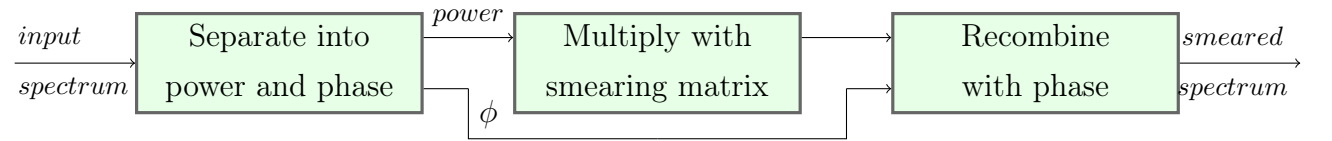


Figure 3: Diagram of the smearing algorithm. Only the power spectrum is processed so that the input and output spectrum have the same phase in order to preserve the waveform and limit noise [2].

3 *MATLAB* implementation

3.1 Functions description

3.1.1 `A_s = calc_smear_matrix(fs, N, b)`

Calculates the smearing matrix A_s according to Baer and Moore article [2]:

$$A_s = A_n^{-1} \cdot A_w$$

where A_n and A_w are normal and broadened auditory filters calculated with `calc_audit_filt.m`.

Input arguments

- f_s : Sampling frequency
- N : Length of the signal. When block processing is used, this corresponds to a frame length.
- b : Broadening factor: $b > 1$, $b = 1$ does not affect the spectrum.

Output arguments

- A_s : Smearing matrix of size $(\frac{N}{2} \times \frac{N}{2})$ to be used in `smearing.m`.

3.1.2 `A = calc_audit_filt(fs, N, b)`

Calculates the broadened auditory filter A according to Baer and Moore article [2]. See section 2 page 1 for more details.

$$A(n, i) = \frac{1}{ERB} = \left(1 + \frac{pg}{b}\right) e^{-\frac{pg}{b}}$$

Input arguments

- f_s : Sampling frequency
- N : Length of the signal. When block processing is used, this corresponds to a frame length.
- b : Broadening factor: $b > 1$, $b = 1$ does not affect the spectrum. This argument is optional, if not given, the default value is 1 (no broadening).

Output arguments

- A : Auditory filter matrix of size $(\frac{N}{2} \times \frac{N}{2})$ to be used in `calc_smear_matrix.m`.

3.1.3 `Y = smearing(X, A_s)`

Smears the power spectrum X (column vector) using the smearing matrix calculated with `calc_smear_matrix.m`. It assumes a real-valued time-domain signal and therefore that spectral component for negative frequencies are the complex conjugate of the positive ones.

Input arguments

- X : Input power spectrum to be smeared given as a column vector of length N .
- A_s : Smearing matrix calculated with `calc_smear_matrix.m` whose dimensions are $(\frac{N}{2} \times \frac{N}{2})$.¹

Output arguments

- Y : Smeared power spectrum given as a column vector of length N .

3.2 Test scripts

Several scripts were used to test different functionalities of the program, separately or combined.

3.2.1 `test_frequency_smearing.m`

This script tests the frequency smearing algorithm for one single frame. To keep computation time low, the signal's length should not exceed 10000 samples. The second part test computing time by calculating a smearing matrix a multiplying it with a spectrum for a large number of growing different frame lengths. The time required to do these operations is measured with the function `tic toc` and plotted versus the frame length. First a test signal is synthesized.

```

1 fs = 16000;          % sampling frequency
2 T = 0.5;             % signal duration
3 f = 1000;            % signal fundammental frequency
4 t = [0:1/fs:T-1/fs]; % time vector
5 x = (sin(2*pi*f*t))' + 0.75*(sin(2*pi*2*f*t))' ...
6   + 0.5*(sin(2*pi*4*f*t))';

```

Then smearing matrix is calculated and the power spectrum is smeared.

```

1 % calculate smearing matrix and output power spectrum
2 A_s = calc_smear_matrix(fs, l, b);
3 Y = smearing(abs(X), A_s);

```

¹The size is $(\frac{N}{2} \times \frac{N}{2})$ and not $(N \times N)$ because only the positive frequencies have to be processed. The negative ones are their complex conjugate.

IFFT is used to recover a time-domain signal than can be analysed as shown in figures 4.

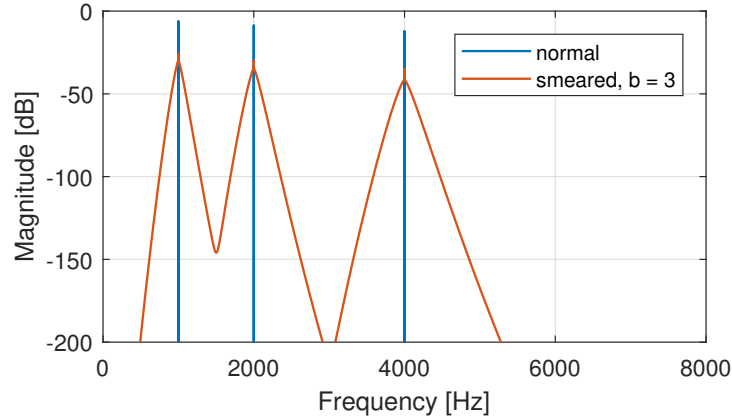


Figure 4: Spectrum of the smeared signal. It is clear that each spikes of the input spectrum is smeared with the shape of the auditory filter. One can notice that the effect is more important at higher frequencies as auditory filter is wider.

3.2.2 frequency_smearing.m

This script loads an audio file and apply frequency smearing to its spectrum. Input signal is processed frame by frame and reconstructed using overlap-and-add method. The overlap is 50% what should achieve perfect recombination with two square-root of hann windows. Note that the 'periodic' argument is used for correct overlap-and-add. Frequency domain output is shown in figure 5, and one can see that overlap-and-add method adds some noise in the spectrum. It will be discussed more in section 3.3.




Figure 5: Spectrum of the smeared signal. The general shape is similar to the one in figure 4, however some periodic spikes appear due to block processing. This may come from small discontinuities in the time domain signal. It will also be shown in the next section that a more complex signal is less affected than single tones.

3.3 Comparison with Bear and Moore results [2]

To replicate the test presented in FIG.2 and FIG.4 of Baer and Moore's article, the following procedure was implemented:

- A test signal containing the word "now" (the vowel /æ/ was used in the article) was sampled at $f_s = 16kHz$ and low-pass filtered at $7kHz$ with the function `FIR_eq`.
- Frame length is 128 samples and a Hamming window is used.
- Each frame is zero-padded with 64 zeros at both ends to increase spectral resolution.

Figures 6 and 7 are replicas of FIG.2 and FIG.4 of the article. It is interesting to notice that the overlap-and-add brings some details back in the spectrum and thus potentially decrease the effect of the spectrum smearing.

In general those results are similar to the ones presented in [2]. Each frame is correctly smoothed and the effect of overlap-and-add described above is also present. This suggests that the implemented algorithm is operating correctly. This will be more discussed in section 5.

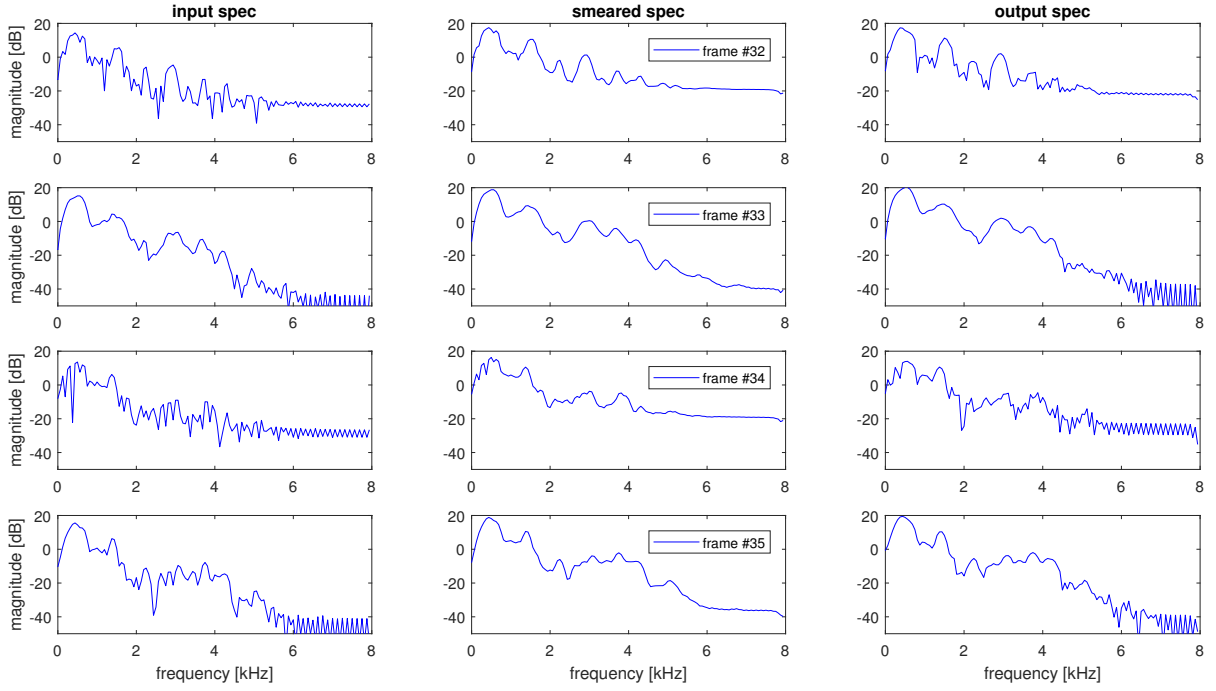


Figure 6: Spectrum of a few frames before frequency smearing (left column), after frequency smearing (middle column) and after overlap-and-add (right column).




Figure 7: Time-domain representation of the zero-padded frames before and after frequency smearing. The black signal is the input signal (on the left) and the overlapped signal (on the right).

3.4 Scripts and functions used for C-code generation

As the smearing matrix is unique for given parameters, it is first calculated in MATLAB and hard coded in the memory of the TEENSY by simply defining arrays. As memory usage is crucial in an embedded system, the row-indexed sparse storage method was used to store the matrix's coefficient in a more judicious way.

3.4.1 Row-indexed sparse storage

The smearing matrix as calculated with `calc_smar_matrix.m` reach very considerable size, depending of the frame length used. For instance, if $N = 256$ samples the smearing matrix has $\left(\frac{N}{2}\right)^2 = 128^2$ elements which correspond to 65.5kB of memory if stored in floating point format. This occupies already a quarter of the TEENSY 3.6 RAM memory (256kB), leaving less space for buffers and other variables. However due to the exponential shape of the auditory filters, most of the matrix's elements are very small relative to the peak resulting in a band diagonal shaped sparse matrix as shown in figure 8. It is then possible to use the *row-indexed sparse storage mode* [1] that only stores the non-zero elements. This compression is done with the script `compress_matrix.m`. The gain of memory depends on the smearing factor b and the threshold under which an element is assumed to be null.

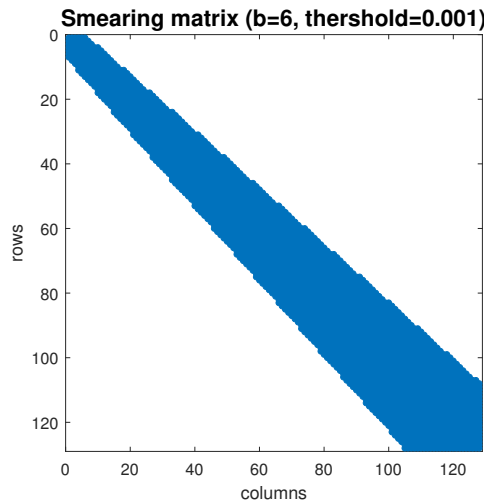


Figure 8: Non-zero elements for a 128 by 128 smearing matrix (all values smaller than the threshold are set to 0). For these particular parameters, row-indexed sparse storage saves 42.2% of memory compared to storing the whole matrix.

3.4.2 generate_smar_matrix.m

This scripts prints the coefficient in a text file using either row-indexed sparse storage or normal storage using two-dimensional arrays. The content of this file can then

be copy-pasted in the C-code (in `smear_mat.h`) to declare the arrays containing the smearing coefficient corresponding to the desired parameters. For instance the following parameters:

```

1 fs = 16e3;      % sampling frequency
2 N  = 128;      % frame length
3 b  = 6;        % smearing coefficient
4 tol = 1e-3;     % thershold used for compressed storage
5 compressed = 1; % 1 -> row-indexed sparse storage
6          % 0 -> classical two-dimensional array

```

will output this line of code for normal storage:

```

1 float smear_mat_b6[2][2] = {{1.000000, 0.000000},
2 {0.000000, 1.000000}};

```

and those two lines for row-indexed sparse storage:

```

1 unsigned int ija_b6[3] = {3, 3, 3};
2 float sa_b6[2] = {1.000000, 1.000000};

```

In this particular example the length is so small that the classical storage is preferable, nonetheless it does not correspond to an real situation.

4 Implementation on *TEENSY* board

4.1 Hardware description

Platform used is a TEENSY 3.6 board¹, that uses an ARM Cortex-M4F 32 bit micro-controller, clocked at 180MHz.

It is combined to the TEENSY audio shield² that provides a high quality 16 bits audio interface. In addition to this, peripheral components as a small electret microphone³ or a jack cable and push-button are used. Figure 9 illustrates how they should be connected to the board.




Figure 9: Wiring schematic to use the microphone. A mono jack cable can also be used by connecting it between *lineIN* and *gnd*. R should set the maximal current to 10mA.

4.2 Software description

The program can be divided according to the different tasks that need to be executed:

- **Real-time block processing.** This part operates the whole data flow, from fetching data from the audio shield, operating input and output circular buffers and sending the data back to the audio shield. The `Audio.h` library [7] and the `CircularBuffer.h` library [6] are used in this part.
- **FFT and IFFT transforms.** The CMSIS DSP library [5] is used for all transforms functions as well as for vector operations. For compatibility with the TEENSY audio library the version 1.1.0 must be used. Therefore it is not needed to install the full library, one can simply include `arm_math.h` that comes with the installation of TEENSY.

¹<https://www.pjrc.com/store/teensy36.html>, (May 2018)

²https://www.pjrc.com/store/teensy3_audio.html, (May 2018)

³<https://www.adafruit.com/product/1713>, (May 2018)

- **Smearing algorithm.** The smearing matrices are calculated in MATLAB and hard-coded in the memory. Therefore the smearing function is simply a matrix multiplication. However to save space in the limited memory, the *row-index sparse storage method* is used.

4.2.1 Real-time block processing

During the execution, audio data pass through several buffer under different formats as illustrated in figure 10. Most important functions are in the files `circ_buff_util.h/cpp`.




Figure 10: Data flow through the different buffers.

1. Audio signal is sampled by the audio shield and theses data are fetched with the `AudioRecordQueue` block from the Teensy audio library. It gives a pointer on a `QUEUE_LEN` samples array of 16 bits integers.
2. It is converted into a floating point with the function `read_array_from_queue()` and pushed at the head of the input circular buffer.

```

1  if(read_array_from_queue(arrayIn, QUEUE_LEN, &queueIn)) {
2      // copy input signal in arrayIn in blocks
3      // copy elements in buffer
4      for(i=0; i<QUEUE_LEN; i++) {
5          buffIn.push(arrayIn[i]);
6      }
7  }

```

3. A frame of `FFT_LEN` samples is read at the tail of the circular buffer with 50% overlap using `read_array_form_buffer()` and converted into a complex array according to the CMSIS convention: `frameC[] = {real[0], imag[0], real[1], imag[1], ...}` and windowed. It is now ready for FFT and processing.
4. The real part of the processed complex frame is windowed and pushed to the head of the output circular buffer, overlapping the last half of the previous frame using the function `overlap_add()`.
5. Blocks of `QUEUE_LEN` samples are read at the tail of the buffer, converted back to 16 bits integer and outputted to the audio shield's headphone line via the `AudioPlayQueue` block.

4.2.2 Smearing algorithm

As explained in section 2 the smearing consist of calculating the power spectrum, multiplying it with the pre-calculated matrix and converting it back to real and imaginary part. This is done with the functions written in `smearing.h/cpp`. The following code extract illustrate it for classical storage.

```

1  for(i=0; i<N/2; i++) {
2      row = smear_mat_b6[i];
3      arm_dot_prod_f32 (row, spec_power, N/2, &tmp);
4      frame[2*i]      = tmp * cos(spec_phase[i]);          // real part
5      frame[2*i+1]    = tmp * sin(spec_phase[i]);          // imaginary part
6  }
```

The particularity of these files is that not everything is compiled, depending in the constant parameters defined in `global.h`. This is needed because only the necessary smearing matrix is declared in `smear_mat.h`.

Table 1 summarizes elapsed time for execution of different tasks. With the same parameters, the duration of FFT/IFFT and smearing process should not excess **7.2 μ s** to avoid audible cuts in the sound.

task	approximative duration
in-to-out buffer path	200 <i>ns</i>
FFT/IFFT	500 <i>ns</i>
smearing	5.8 μ s
total	6.5 μs

Table 1: Example of processing times for $f_s = 16kHz$ and a frame length of `FFT_LEN = 256`.

4.3 Functions description

4.3.1 read_array_form_queue

Arguments

- `float arrayIn[]`. Pointer to the array to be filled with samples.
- `int array_length`. Length of the array. It must be a multiple of `QUEUE_LEN`. If a smaller size is wanted, `QUEUE_LEN` can be changed in `AudioStream.h`.
- `AudioRecordQueue* queueIn`. Pointer to the input queue that provide data form the audio shield.

Description

This function converts and copies input signal fetched from the input queue into a floating point array. Returns 0 if the dimensions mismatch or if the queue does not have enough samples.

4.3.2 read_frame_from_buffer*Arguments*

- `float frame[]`. Pointer to the array to be filled in with elements from the input circular buffer `buffIn`.
- `int frame_l`. Length of the frame.

Description

This function copies `frame_l` samples from the input circular buffer and leaves the last half of them in the buffer for the next frame. It results in 50% overlap.

4.3.3 overlap_add*Arguments*

- `float frame[]`. Array of elements (real-valued) to be added in the output circular buffer `buffOut`.
- `int array_length`. Length of the frame.

Description

This function pushes the elements of the frame to the head of the output circular buffer, overlapping the last half of the previous frame.

4.3.4 smearing_uncomp and smearing_comp*Arguments*

- `float32_t* frame`. Pointer on the *complex* vector to be smeared.
- `int N`. Number of complex elements in the vector.

Description

Both functions process frequency smearing according to the parameters defined in `global.h`. `smearing_uncomp` uses a classic matrix multiplication and `smearing_comp` uses row-indexed sparse storage multiplication done in `spr sax`.

Note 1: Only one of these two functions is compiled, depending if the symbol `COMPRESSED` is defined.

Note 2: They are written to work with $B = 3$ and $B = 6$. If different options are wanted, new compiling directive must be written with the corresponding smearing matrix written in `smear_mat.h`:

```

1 #elif B == NEWVALUE
2     for(i=0; i<N/2; i++) {
3         row = smear_mat_bNEWVALUE[i];
4         arm_dot_prod_f32 (row, spec_power, N/2, &tmp);
5         frame[2*i]      = tmp * cos(spec_phase[i]);           // real part
6         frame[2*i+1]    = tmp * sin(spec_phase[i]);           // imaginary part
7     }

```

4.3.5 sprsax

Arguments

- `float sa[]` and `unsigned int ija[]`. Arrays containing the non-zero coefficient of the smearing matrix and indexes needed for the row-indexed sparse storage. The first element of `ija` must be equal to `n+1`.
- `float x[]`. Vector to be multiplied.
- `float b[]`. Output vector.
- `unsigned int n`. Length of the vector \mathbf{x} and \mathbf{b} .

Description

This functions process the vector-by-matrix multiplication in equation 7 where \mathbf{A} is stored with row-indexed sparse storage. It is directly adapted from [1].

$$\mathbf{b} = \mathbf{A}\mathbf{x} \quad (7)$$

4.3.6 neg_freq

Arguments

- `float32_t* frame`. Input and output complex spectrum where frequencies are sorted as:

$$[0Hz, \Delta_f, \dots, (f_{max} - \Delta_f), f_{max}, -f_{max}, (-f_{max} + \Delta_f), \dots, -\Delta_f]$$

- `int N`. Number of complex elements in frame.

Description

This function reconstructs spectrum's negative frequencies with complex conjugate of the positive one.

4.3.7 `fft_init`

Arguments

- `arm_cfft_radix2_instance_f32* fftInst`. Pointer to the instance of the floating-point CFFT/CIFFT structure to be initialized.
- `uint16_t fftLen`. Length of the FFT/IFFT in number of complex elements.
- `uint8_t fftFlag`. Flag that selects forward (`fftFlag=0`) or inverse (`fftFlag=1`) transform.
- `uint8_t bitReverseFlag`. Flag that enables (`bitReverseFlag=1`) or disables (`bitReverseFlag=0`) bit reversal of output. It is always enabled in this program.

Description

This function initializes the FFT and IFFT instances using the CMSIS function `arm_cfft_radix2_init_f32`¹ and checks for errors.

Note 1: The CMSIS documentation specifies that the supported length are `fftLen = 16, 64, 256, 1024`. However on one thread of the TEENSY forum² one can read:

"The radix2 ones offer 16, 32, 64, 128, 256, 512, 1024, 2048 sizes. On Cortex-M4 for Q15, the radix4 versions is slightly faster, which is the reason it was used with the audio library."

what would suggest more flexibility towards possible length. It was successfully tested for `FFT_LEN = 128` and `512`.

Note 2: `FFT_LEN` should always be larger than `QUEUE_LEN`. If a smaller value than 128 samples is wanted, `AUDIO_BLOCK_SAMPLES` has to be changed in `AudioStream.h` using multiples of 16 samples.

4.3.8 `set_periph`

Description

This function initializes the SD card, the buttons, the audio shield and set the I2S sampling frequency. Its only purpose is encapsulation to improve code readability. To simplify inclusions and parameters it is directly written in the `.ino` file and not in `util.h/cpp`.

¹https://www.keil.com/pack/doc/CMSIS/DSP/html/group__ComplexFFT.html

²<https://forum.pjrc.com/threads/35277-arm-math-h-and-the-FFT-audio-blocks>

4.3.9 setI2SFreq

Arguments

- `int freq`. Wanted sampling frequency, possible values are (in Hertz): 8000, 11025, 16000, 22050, 32000, 44100, 44117 , 48000, 88200, 88234, 96000, 176400, 176468 and 192000.

Description

This function adapts the audio shield's I2S in- and output sampling frequency. It is copied from the TEENSY forum¹.

Note: The sampling frequency is only adapted for the I2S. The rest of the audio library still works on $44.1kHz$. Therefore, if other modules as the waveform generator are used their frequency must be scaled:

```
1 waveform1.frequency(AUDIO_SAMPLE_RATE_EXACT/FS*freq); // where freq is the desired
    frequency in Hz
```

4.3.10 read_button

Description

This function check if the button connected to `BUTTON1_PIN` has been pushed and (des-)activates frequency smearing. If more buttons are needed some code can easily be added following the same scheme.

4.3.11 create_hann_window and create_sqrthann_window

Arguments

- `float win[]`. Pointer to the window vector to be initialized.
- `int win_l`. Length of the window.

Description

Those two function create a *hann* or \sqrt{hann} window respectively. It uses the `AudioWindowHanning256` defined in `Audio.h`, therefore 256 is the only possible length. More flexibility could be added by defining local variables with window coefficients for other lengths.

¹<https://forum.pjrc.com/threads/38753-Discussion-about-a-simple-way-to-change-the-sample-rate>

5 Measurements on TEENSY

When listening to different audio files processed by the on the embedded platform, one can clearly notice the effect of frequency searing. However a same signal does not sound exactly the same when processed in MATLAB or on the TEENSY. In order to have more objective appreciation of the system's effect the audio shield's line level output was used to recorded some samples through an audio interface.

5.1 Experimental setup

The software does not need to be changed as the audio is as the `queueOut` block automatically sends audio to the headphone output and the line level out.




Figure 11: Schematic of the setup used to record audio from the TEENSY. Recordings were made at a sampling frequency of $128kHz$ with a $24bits$ resolution. A high-pass filter and low-pass filter were added in the recording software to reduce noise under $1kHz$ (mainly due to 50Hz power supply) and above $8kHz$.

In addition to those recordings the smearing algorithm alone can also be tested using the sketch `test_algo.ino`. This latter synthesizes a simple audio signal of finite length and processes it once using the same smearing function and overlap-and-add as in the main program. The in- and output signals are then written into a text file stored on the SD card to be analyzed in MATLAB.

5.2 Results and discussion

The same signal as in figures 4 and 5 in section 3 was used and results of the recordings can be seen on figure 12 for which several points need to be discussed:

- Noise level is higher for the smeared spectrum, this is also noticeable when listening to TEENSY's output. This might also be a reason why frequency smearing has greater effects on speech intelligibility in presence of noise [2].
- Repetitive spikes appear every 125Hz as in MATLAB block processing (figure 5). This frequency correspond to $f = (\frac{0.5 \cdot FFT_{LEN}}{f_s})^{-1}$ what suggests that it is an effect from block processing and incorrect reconstruction in overlap-and-add. While it is quite important when simple tones are used as test signals, more complex audio signals are less affected by this effect.
- One can also notice some peaks at higher frequencies ($5kHz$, $6kHz$ and $7kHz$). At this point no plausible explanation has been found and more recordings should be done to determine if it is really a problem in the program or some perturbation introduced by the setup.
- Finally one can be surprise to see that smearing is barely noticeable despite a rather high coefficient b . A possible explanation is that it is effectively masked by the too high noise level. Indeed, as shown in figure 13, the smeared side of each frequency peak has a level approximately $45dB$ lower than the peak itself what is comparable to the signal-to-noise ratio of the recorded smeared audio. To confirm it, better recordings has to be done and it would be judicious to use more complex audio signal where the energy is more spread across the spectrum. This was tried but the high noise level make the results useless.




Figure 12: Spectrum of the signal recorded using the setup shown in figure 11. Frequency resolution has been reduced on purpose to achieved better readability. The three frequencies of the signal ($1kHz$, $2kHz$ and $4kHz$) are clearly identifiable but more components seems to be present, especially for the smeared signal. Repetitive spikes appear every $125Hz$ as in MATLAB block processing.

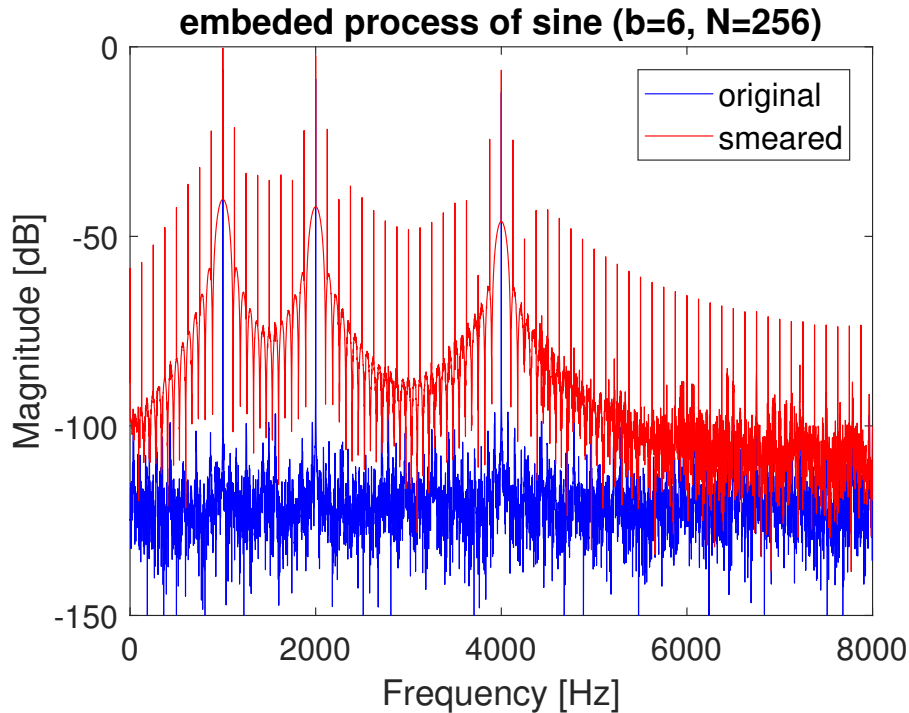


Figure 13: Spectrum of the signal synthesized, processed and recorded using the second method described above. Its advantage is that it tests only the smearing algorithm and eliminates possible influence of buffers and especially audio path outside of the TEENSY. Spikes every $125Hz$ are also present and frequency smearing is clearly seen. However, the smeared part of each frequency peak has a level approximately $45dB$ lower than the the peak itself.

6 Limitations, improvements and future work

In addition to the issues discussed on the previous section, some features need to be improved or added to the program. Important points are, among others, available processing time, memory usage, user interface and flexibility.

6.1 Processing time

The biggest challenge regarding processing time is the ability to run the program in real-time with a continuous output. As presented in section 4.2.2 on page 13, transforming the frame and processing frequency smearing takes approximately $6.3\mu s$ for a frame length of 256 samples using row-indexed storage. This time would obviously be decreased by reducing the frame length, however this leaves less time for processing between two consecutive frames. This is why the sampling frequency is reduced from $44.1kHz$ to $16kHz$, with the consequence to lose the high frequencies of the processed audio. Nevertheless one could discuss the importance of this inconvenient as hearing impairment also induces losses of high frequencies, even to a lower level.

6.2 Memory usage, operation flexibility and characterization

With *only* 256 kbytes of RAM, storing huge smearing matrices and long buffers is not possible on the embedded platform. This is why the approach of compiling only the useful code using directive was taken. Unfortunately this reduces a lot the possibilities of changing parameters as smearing coefficient or even frame length while the program is executing. Moreover the implementation of row-indexed sparse storage would allow to have a few different matrices stored at the same time. Implementing different coefficients and the ability to switch between them on-line would be a big improvement but it implies a important restructuring of the code in `smear_mat.h` and the smearing function.

It would also be very interesting to be able to vary the frame length, even if it is not on-line, in order to investigate the performance of the algorithm for different frequency resolution (the tests on MATLAB showed that the smearing is more noticeable with high frequency resolution, thus longer frames). To do so, some more hann window have to be implemented as explained in section 4.3.11.

6.3 Hardware improvement

Following the objective of a demonstration kit, a better and more "finished" hardware implementation is needed, including adjustable volume and/or gain, buttons to activate smearing and change smearing coefficient. This last point could even be done using wireless command or some timer so that it is not directly commanded by the user. An ideal implementation would be integrated in some headphones.

Adding a low-pass filter with adjustable cutoff frequency would approach hearing impairment more realistically as would stereo processing.

6.4 Algorithm improvement

In their article [2], Baer and Moore point that asymmetrical broadening around the center frequency has a stronger effect on speech intelligibility. It is also closer to the actual broadening that happens in hearing impaired subject. To implement this, the MATLAB functions have to be modified using different b factor for each half of the spectrum.

6.5 Future work

The points described above can be summarized in a task list:

- Possibility to change smearing coefficient in-line.
- Implement more different window lengths to investigate effect of frequency resolution.
- Operate more and better recordings to have precise and rigorous results about the operation on the embedded platform.
- Add adjustable low-pass filter.
- Implement asymmetrical broadening.
- Implement stereo processing.
- Improve hardware packaging and user interface to have an operational demonstration kit.

References

- [1] William H. Press et al. *Numerical Recipes in C: The Art of Scientific Computing, Second Edition*. English. 2 edition. Cambridge ; New York: Cambridge University Press, Oct. 1992. ISBN: 978-0-521-43108-8.
- [2] Thomas Baer and Brian C. J. Moore. “Effects of Spectral Smearing on the Intelligibility of Sentences in Noise”. en. In: *The Journal of the Acoustical Society of America* 94.3 (Sept. 1993), pp. 1229–1241. ISSN: 0001-4966. DOI: 10.1121/1.408176.
- [3] Dirk Mauler and Rainer Martin. “A Low Delay, Variable Resolution, Perfect Reconstruction Spectral Analysis-Synthesis System for Speech Enhancement”. en. In: (2007), p. 5.
- [4] Andreas Spanias, Ted Painter, and Venkatraman Atti. *Audio Signal Processing and Coding*. en. Hoboken, NJ, USA: John Wiley & Sons, Inc., Jan. 2007. ISBN: 978-0-470-04197-0 978-0-471-79147-8. DOI: 10.1002/0470041978.
- [5] *CMSIS DSP Software Library*. <https://www.keil.com/pack/doc/CMSIS/DSP/html/index.html>. May 2018.
- [6] Roberto Lo Giacco. *CircularBuffer: Arduino Circular Buffer Library*. <https://github.com/rlogiacco/CircularBuffer>. Apr. 2018.
- [7] *Teensy Audio Library, High Quality Sound Processing in Arduino Sketches on Teensy 3.1*. https://www.pjrc.com/teensy/td_libs_Audio.html. May 2018.